

MLP-Mixer

An all-MLP Architecture for Vision


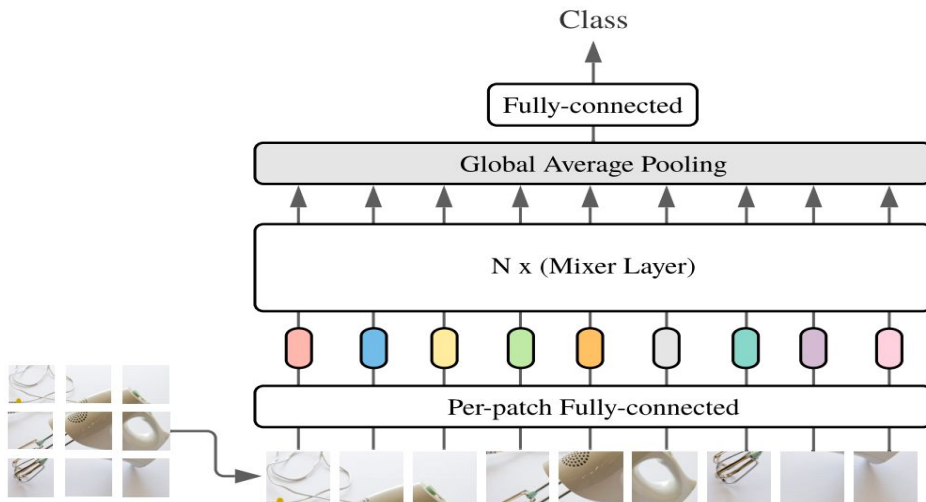
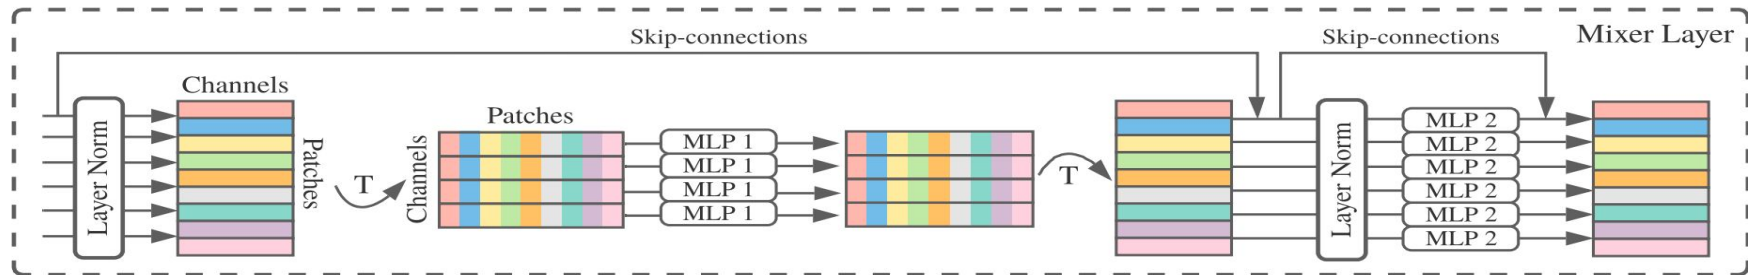
Sebenele Thwala

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

IDEA

- Alternative and Simpler approach to vision networks
- Show that CNNs and Attention although sufficient for good performance, are not necessary

ARCHITECTURE



MLP BLOCK

```
class MLP(nn.Module):  
    """  
    dim is the input dimension in the MLP, and hidden_dim is the dimensions of the hidden layers  
    """  
    def __init__(self, dim, hidden_dim):  
  
        super().__init__()  
        self.lin1 = nn.Linear(dim, hidden_dim)  
        self.gelu = nn.GELU()  
        self.lin2 = nn.Linear(hidden_dim, dim)  
  
    def forward(self, x):  
        x = self.lin1(x)  
        x = self.gelu(x)  
        x = self.lin2(x)  
  
        return x
```

Mixer Layer

```
class MixerLayer(nn.Module):
    def __init__(self, n_patches, dim, token_dim, channel_dim ):

        super().__init__()

        self.norm1 = nn.LayerNorm(dim)
        self.token_mix = MLP(n_patches, token_dim)
        self.channel_mix = MLP(dim, channel_dim)
        self.norm2 = nn.LayerNorm(dim)

    def forward(self, x):

        y = self.norm1(x)
        y = self.token_mix(torch.transpose(y, 1, 2))
        y = torch.transpose(y, 1, 2)
        x = x + y
        y = self.norm2(x)
        out = x + self.channel_mix(y)

        return out
```

```

class MLPMixer(nn.Module):
    def __init__(self, img_sz, patch_sz, dim, token_dim, channel_dim, n_classes, n_blocks):
        super().__init__()

        n_patches = (img_sz//patch_sz)**2

        #to extract patch from input images
        self.patch_kernel = nn.Conv2d(3, dim, kernel_size = patch_sz, stride= patch_sz)

        #our N mixer blocks
        self.blocks = nn.ModuleList(
            [
                MixerLayer(
                    n_patches = n_patches,
                    dim = dim,
                    token_dim = token_dim,
                    channel_dim = channel_dim,
                )
                for _ in range(n_blocks)
            ]
        )
        self.pre_head_norm = nn.LayerNorm(dim)
        self.head_classifier = nn.Linear(dim, n_classes)

    def forward(self, x):

        x = self.patch_kernel(x) #input to mixerlayers is patch in tensor form not image
        x = einops.rearrange(x, "n c h w -> n (h w) c")

        for mixer_block in self.blocks:

            x = mixer_block(x)
            x = self.pre_head_norm(x)
            x = x.mean(dim = 1) #avarage pooling
            y = self.head_classifier(x)

        return y

```

EXPERIMENTS

epoch=0, train=1: 100%		782/782	[07:28<00:00,	1.74it/s,	acc=0.3285,	loss=1.8335]
epoch=0, train=0: 100%		157/157	[00:44<00:00,	3.50it/s,	acc=0.4216,	loss=1.6007]
epoch=1, train=1: 100%		782/782	[07:28<00:00,	1.74it/s,	acc=0.4433,	loss=1.5298]
epoch=1, train=0: 100%		157/157	[00:44<00:00,	3.50it/s,	acc=0.4688,	loss=1.4704]
epoch=2, train=1: 100%		782/782	[07:34<00:00,	1.72it/s,	acc=0.4775,	loss=1.4386]
epoch=2, train=0: 100%		157/157	[00:45<00:00,	3.49it/s,	acc=0.4873,	loss=1.4345]
epoch=3, train=1: 100%		782/782	[07:29<00:00,	1.74it/s,	acc=0.5002,	loss=1.3820]
epoch=3, train=0: 100%		157/157	[00:45<00:00,	3.47it/s,	acc=0.5071,	loss=1.3591]
epoch=4, train=1: 100%		782/782	[07:29<00:00,	1.74it/s,	acc=0.5188,	loss=1.3348]
epoch=4, train=0: 100%		157/157	[00:44<00:00,	3.50it/s,	acc=0.5228,	loss=1.3403]
epoch=5, train=1: 100%		782/782	[07:30<00:00,	1.74it/s,	acc=0.5263,	loss=1.3049]
epoch=5, train=0: 100%		157/157	[00:45<00:00,	3.46it/s,	acc=0.5201,	loss=1.3165]
epoch=6, train=1: 100%		782/782	[07:29<00:00,	1.74it/s,	acc=0.5412,	loss=1.2732]
epoch=6, train=0: 100%		157/157	[00:44<00:00,	3.50it/s,	acc=0.5390,	loss=1.2793]
epoch=7, train=1: 100%		782/782	[07:29<00:00,	1.74it/s,	acc=0.5552,	loss=1.2410]
epoch=7, train=0: 100%		157/157	[00:45<00:00,	3.48it/s,	acc=0.5393,	loss=1.2745]
epoch=8, train=1: 100%		782/782	[07:28<00:00,	1.74it/s,	acc=0.5653,	loss=1.2115]
epoch=8, train=0: 100%		157/157	[00:45<00:00,	3.49it/s,	acc=0.5535,	loss=1.2384]
epoch=9, train=1: 100%		782/782	[07:33<00:00,	1.72it/s,	acc=0.5779,	loss=1.1851]
epoch=9, train=0: 100%		157/157	[00:45<00:00,	3.43it/s,	acc=0.5562,	loss=1.2237]

Comparison

upstream	model	dataset	accuracy	wall_clock_time
ImageNet	Mixer-B/16	cifar10	96.72%	3.0h
ImageNet	Mixer-L/16	cifar10	96.59%	3.0h
ImageNet-21k	Mixer-B/16	cifar10	96.82%	9.6h
ImageNet-21k	Mixer-L/16	cifar10	98.34%	10.0h

CONCLUSION